

Class 7 Machine Learning 1

Melissa (PID: A19149673)

Table of contents

Background	1
K-means clustering	1
Hierarchical Clustering	6
Principal Component Analysis (PCA)	8
Analysis of UK food data	8
Data Import	8
Tidy the data	9
Exploratory analysis	9
PCA	9
Complete the following code to find out how many rows and columns are in x? . . .	9
PCA to the rescue	16

Background

Today we will explore some core machine learning methods that are very popular in bioinformatics. These include **clustering** and **dimensionality reduction**.


K-means clustering

The main function in “base” R for K-means clustering is called `kmeans()`


Before we go too deep let’s make up some “simple” data that we can cluster and know if we are getting a good answer or not. To do this we can use `rnorm()` function:

```
hist( rnorm (10000, mean = 3))
```

Histogram of rnorm(10000, mean = 3)



```
x <- c( rnorm(30, -3), rnorm(30, +3))  
z <- cbind(x=x, y = rev(x))  
plot(z)
```



```
p <- 1:5  
rev(p)
```

[1] 5 4 3 2 1

Now we can run `kmeans()` on this input `z` and see what the results look like.

```
km <- kmeans(z, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.912375	3.259432
2	3.259432	-2.912375

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 64.88041 64.88041
(between_SS / total_SS = 89.8 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
attributes(km)
```

\$names

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

\$class

```
[1] "kmeans"
```

Q. How many points are in each cluster?

```
km$size
```

```
[1] 30 30
```

Q. What “component of your result object details cluster assignment/membership”?

```
km$cluster
```

[illegible]


Q. What “component of your result object details cluster center”?

```
km$centers
```

	x	y
1	-2.912375	3.259432
2	3.259432	-2.912375


Q. Plot $\mathbf{z}$ colored by the kmeans cluster assignment and add cluster centers as blue points.

```
plot(z, col=km$cluster)
points(km$centers, col = "blue", pch = 15)
```



Q. Run a K-means clustering and plot the results asking for 4 clusters? (k=4)?

```
km4 <- kmeans(z, centers = 4)
plot(z, col=km4$cluster)
points(km4$centers, col="black", pch = 15)
```




N.B. You need to tell K-means the number of clusters (i.e. set `centers=2`)!!

One approach is to try different values for `centers` and then pick the best...

```
ans <- NULL
for(i in 1:10) {
  km <- kmeans(z, centers=i)
  ans <- c(ans, km$tot.withinss)
}

plot(ans, typ="o", xlab = "Number of clusters",
      ylab = "Total Sum of Squares Distance")
```



Hierarchical Clustering

The main function in “base” R for Hierarchical Clustering is called `hclust()`

This function does not take your “raw” data for clustering. You must first build a “distance matrix” from your data and pass this as input to `hclust()`

```
d <- dist(z)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance         : euclidean
Number of objects: 60
```

There is a bespoke `plot()` method for `hclust()` result objects.

```
plot(hc)
abline(h = 3, col = "red")
```




Once we have our `hclust` object (our “tree” of “cluster dendrogram”) we can “*cut*” the tree to reveal the clustering pattern.

```
cutree(hc, k=4)
```

[illegible]

Q. Make a plot of $\mathbf{z}$ with your hclust results (i.e. colored by cluster membership)

```
grps <- cutree(hc, k = 4)
plot(z, col= grps, pch =16)
```



Principal Component Analysis (PCA)

PCA is a dimensionality reduction method that is popular for revealing patterns in complex datasets.

Analysis of UK food data

Let's look at some data on the eating habits of folks from the UK to see if there are patterns and trends that have some regions being distinct from others.

Data Import

The data is made available in CSV format so we can use `read.csv()` function for import to R:

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x
```


	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139
7	Fresh_potatoes	720	874	566	1033
8	Fresh_Veg	253	265	171	143
9	Other_Veg	488	570	418	355
10	Processed_potatoes	198	203	220	187
11	Processed_Veg	360	365	337	334
12	Fresh_fruit	1102	1137	957	674
13	Cereals	1472	1582	1462	1494
14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

Tidy the data

Fix anything that went wrong with the data import.

Exploratory analysis

Make some plots to help make sense of obvious trends...

PCA

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

Complete the following code to find out how many rows and columns are in x?

```
dim(x)
```

```
[1] 17  5
```

##Preview the first 6 rows

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
rownames(x) <- x[,1] # This keeps only column 1  
x <- x[,-1] #This keeps everything except column 1  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```


There are 17 rows and 4 columns after fixing the row names.

```
x <- read.csv(url, row.names=1)  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Setting row names during import is more robust since it avoids extra steps and prevents error later.

Q3. Changing what optional argument in the above `barplot()` function results in the following plot?

```
cols <-rainbow(nrow(x))  
barplot(as.matrix(x), col=cols)
```



```
dim(x)
```

```
[1] 17  4
```

```
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
library(tidyr)
```


```
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
               names_to = "Country",
               values_to = "Consumption")
```

```
dim(x_long)
```

```
[1] 68  3
```

```
library(ggplot2)
```


```
ggplot(x_long) +  
  aes(x = Country, y = Consumption, fill = Food) +  
  geom_col(position = "dodge") +  
  theme_bw()
```



The change comes from changing the argument that controls how the bars are arranged. Also determines whether the barplot shows separate bars for each category as well as combined stack bars.

Q4: Changing what optional argument in the above `ggplot()` code results in a stacked barplot figure?


```
ggplot(x_long) +  
  aes(x = Country, y = Consumption, fill = Food) +  
  geom_col(position = "stack") +  
  theme_bw()
```




Switching the argument changes the plot from comparing categories separately to show their combined totals within each group.

Q5: We can use the `pairs()` function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap(as.matrix(x))
```



The pairs() plot shows all pairwise comparison between the countries. This helps visually show how similar and different countries are based on their food consumption across the categories.

Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Key-point: Even relatively small datasets can prove challenging to interpret.

PCA to the rescue

`t(x)`

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England	720	253	488		198	
Wales	874	265	570		203	
Scotland	566	171	418		220	
N.Ireland	1033	143	355		187	
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England	360	1102	1472	57	1374	
Wales	365	1137	1582	73	1256	
Scotland	337	957	1462	53	1572	
N.Ireland	334	674	1494	47	1506	
	Alcoholic_drinks	Confectionery				
England	375	54				
Wales	475	64				
Scotland	458	62				
N.Ireland	135	41				

The main function in “base” R for PCA is called `prcomp()`. This function wants the observations to be rows and the variables to be columns.

```
pca <- prcomp( t(x) )
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

```
$names
```

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
```

```
[1] "prcomp"
```

The main result figure from this analysis is called “PC score plot” or “ordination plot” “PC plot” or PC1 vs PC2 plot”.

This plot show samples (in this case countries) relate to each other along our new PC axis.

This is our new “reduced-dimensional space”. In this case 2 dimensions, PC1 and PC2, that capture most of the variance to the original 17 dimensional data-set.

```
pca$x
```


	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

England, Wales and Scotland cluster together to show similar diets while Northern Ireland is pretty distinct. However, the specific food differences require further analysis such as the PCA loading's to be able to identify.

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

ggplot(pca$x) +
  aes(x = PC1, y = PC2, label = rownames(pca$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.


```
df <- as.data.frame(pca$x)
df$Country <- rownames(df)

ggplot(df) +
  aes(x = PC1, y = PC2, label = Country, color = Country) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
```

```

scale_color_manual(values = c(
  "England" = "orange",
  "Wales" = "red",
  "Scotland" = "blue",
  "N.Ireland" = "darkgreen"
)) +
xlim(-270, 500) +
xlab("PC1") +
ylab("PC2") +
theme_bw()

```



```

v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v

```

```
[1] 67 29 4 0
```

```

z <- summary(pca)
z$importance


```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14

Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00


```
# Create scree plot with ggplot
variance_df <- data.frame(
  PC = factor(paste0("PC", 1:length(v)), levels = paste0("PC", 1:length(v))),
  Variance = v
)

ggplot(variance_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  xlab("Principal Component") +
  ylab("Percent Variation") +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 0))
```




```
ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
```

```
ylab("") +
theme_bw() +
theme(axis.text.y = element_text(size = 9))
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
ggplot(pca$rotation) +
  aes(x = PC2,
       y = reorder(rownames(pca$rotation), PC2)) +
  geom_col(fill = "steelblue") +
  xlab("PC2 Loading Score") +
  ylab("") +
  theme_bw()
```



The food groups that feature most prominently in PC2 are those with the largest positive and negative loading scores such as fish and fats_and_oils. These variables contribute strongly to the variation captured by PC2. PC2 also represents a secondary dietary pattern, as this distinguishes countries based on differences in specific types of food consumption rather than overall intake. Moreover, PC1 explains about 67% of the variance while PC2 explains about 29%.